

30.08.2024

# Statistical Methods in AI (CS7.403)

Lecture-8: Clustering (k-means, Gaussian Mixture Models)

Ravi Kiran ([ravi.kiran@iiit.ac.in](mailto:ravi.kiran@iiit.ac.in))

<https://ravika.github.io>



Center for Visual Information Technology (CVIT)

IIIT Hyderabad

Repeat until convergence: {

For every  $i$ , set

$$c^{(i)} := \arg \min_j \|x^{(i)} - \mu_j\|^2.$$

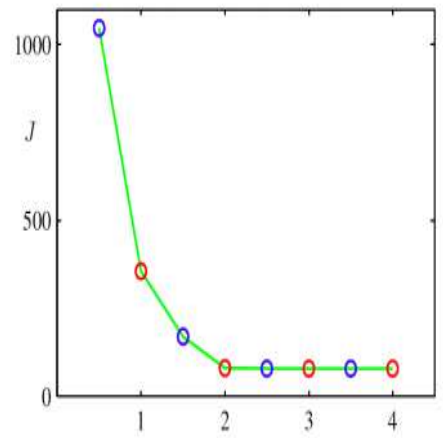
For each  $j$ , set

$$\mu_j := \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}.$$

}

$$\sum_{k=1}^K \sum_{i=1}^{n_k} \|x_{ki} - \mu_k\|^2$$

- Whenever an assignment is changed, the sum squared distances  $J$  of data points from their assigned cluster centers is reduced.
- Whenever a cluster center is moved,  $J$  is reduced.
- **Test for convergence:** If the assignments do not change in the assignment step, we have converged (to at least a local minimum).



- K-means cost function after each E step (blue) and M step (red). The algorithm has converged after the third M step

- The objective  $J$  is non-convex (so coordinate descent on  $J$  is not guaranteed to converge to the global minimum)



1. Initialize **cluster centroids**  $\mu_1, \mu_2, \dots, \mu_k \in \mathbb{R}^n$  randomly.
2. Repeat until convergence: {

For every  $i$ , set

$$c^{(i)} := \arg \min_j \|x^{(i)} - \mu_j\|^2.$$

For each  $j$ , set

$$\mu_j := \frac{\sum_{i=1}^m 1\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m 1\{c^{(i)} = j\}}.$$

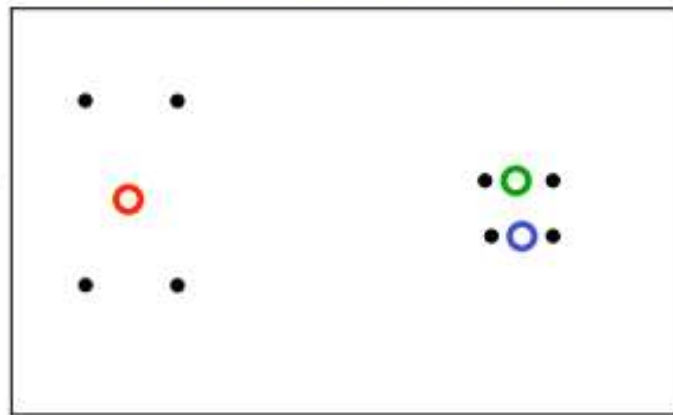
}

$$se = \sum_{k=1}^K \sum_{i=1}^{n_k} \|x_{ki} - \mu_k\|^2$$

- The objective  $J$  is non-convex (so coordinate descent on  $J$  is not guaranteed to converge to the global minimum)
- There is nothing to prevent k-means getting stuck at local minima.



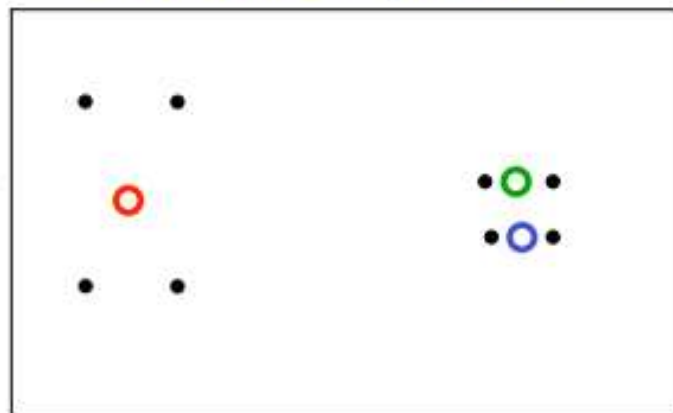
A bad local optimum



- The objective  $J$  is non-convex (so coordinate descent on  $J$  is not guaranteed to converge to the global minimum)
- There is nothing to prevent k-means getting stuck at local minima.
- We could try many random starting points

$$\sum_{k=1}^K \sum_{i=1}^{n_k} \|x_{ki} - \mu_k\|^2$$

A bad local optimum

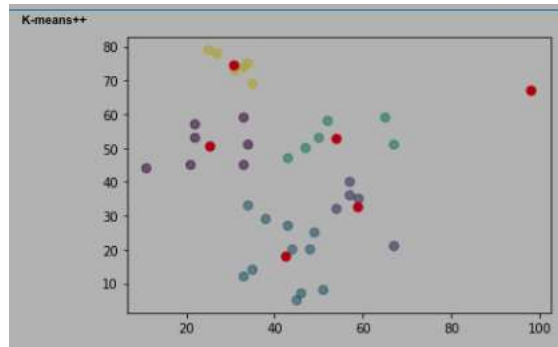
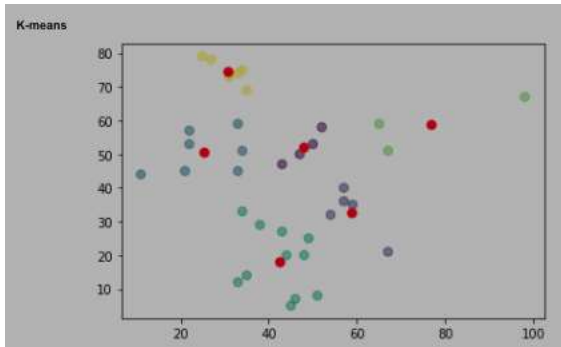


# K-means++: Improving K-means initialization

- Common way to improve k-means - smart initialization!
- General idea - try to get good coverage of the data.
- k-means++ algorithm:
  1. Pick the first center randomly
  2. For all points  $\mathbf{x}^{(n)}$  set  $d^{(n)}$  to be the distance to closest center.
  3. Pick the new center to be at  $\mathbf{x}^{(n)}$  with probability proportional to  $d^{(n)2}$
  4. Repeat steps 2+3 until you have k centers

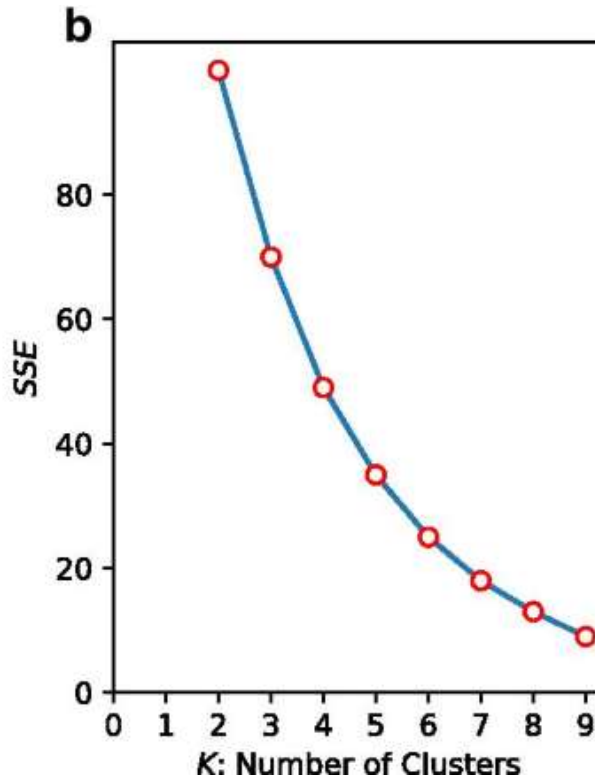
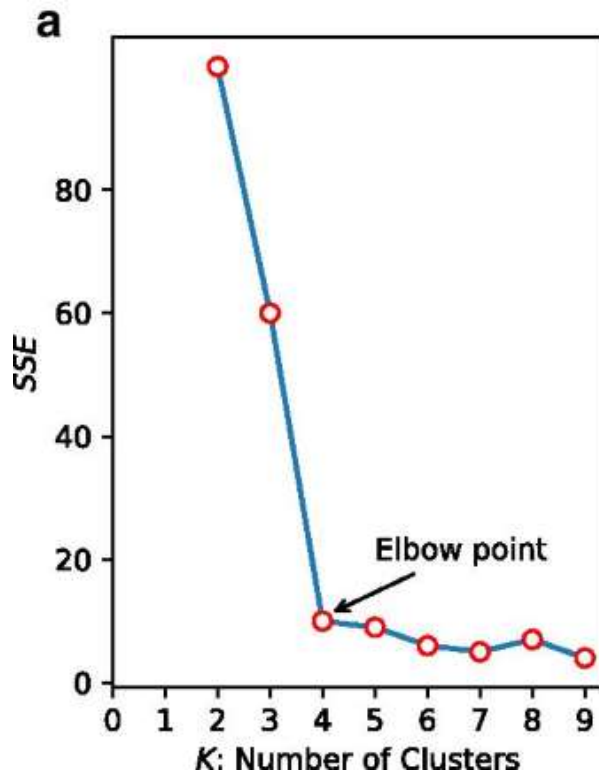
# K-means++: Improving K-means initialization

- Common way to improve k-means - smart initialization!
- General idea - try to get good coverage of the data.
- k-means++ algorithm:
  1. Pick the first center randomly
  2. For all points  $\mathbf{x}^{(n)}$  set  $d^{(n)}$  to be the distance to closest center.
  3. Pick the new center to be at  $\mathbf{x}^{(n)}$  with probability proportional to  $d^{(n)2}$
  4. Repeat steps 2+3 until you have k centers



# How to choose k ?

$$se = \sum_{k=1}^K \sum_{i=1}^{n_k} \|x_{ki} - \mu_k\|^2$$





# Regularization

- Penalize “overly” large or “overly” small clusters

# K-mediods

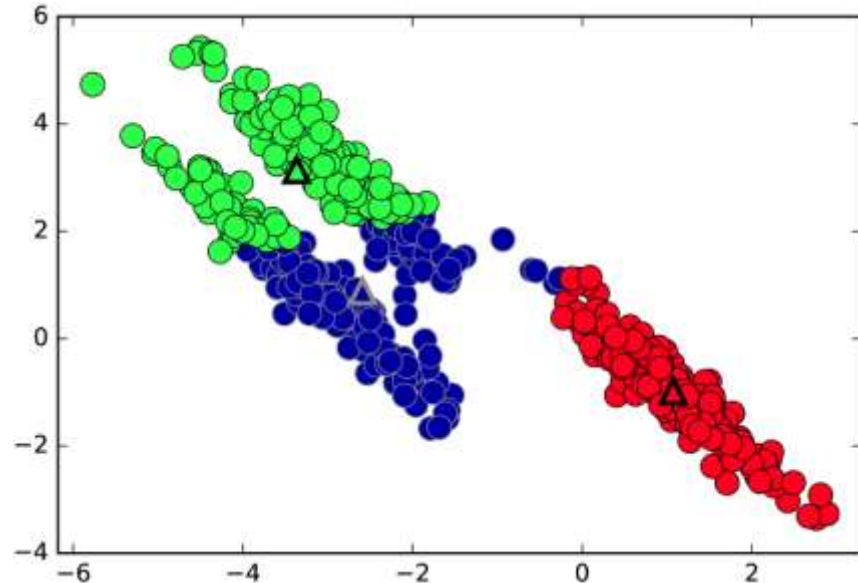
- Squared Euclidean distance loss function of K-means not robust.

- Use L1 loss function  $J = \sum_{i=1}^n \sum_{k=1}^K r_{ik} \|x_i - \mu_k\|_1$  instead of squared Euclidean distance.

- Use an iterative procedure as before.
  - Prototype is the median of the points assigned to a cluster.

# K-means: Additional issues

- 'Hard' assignments
- Euclidean  $\rightarrow$  Favours 'spherical' clusters of equal 'contribution'
- Sensitive to initialization
- Sensitive to outliers



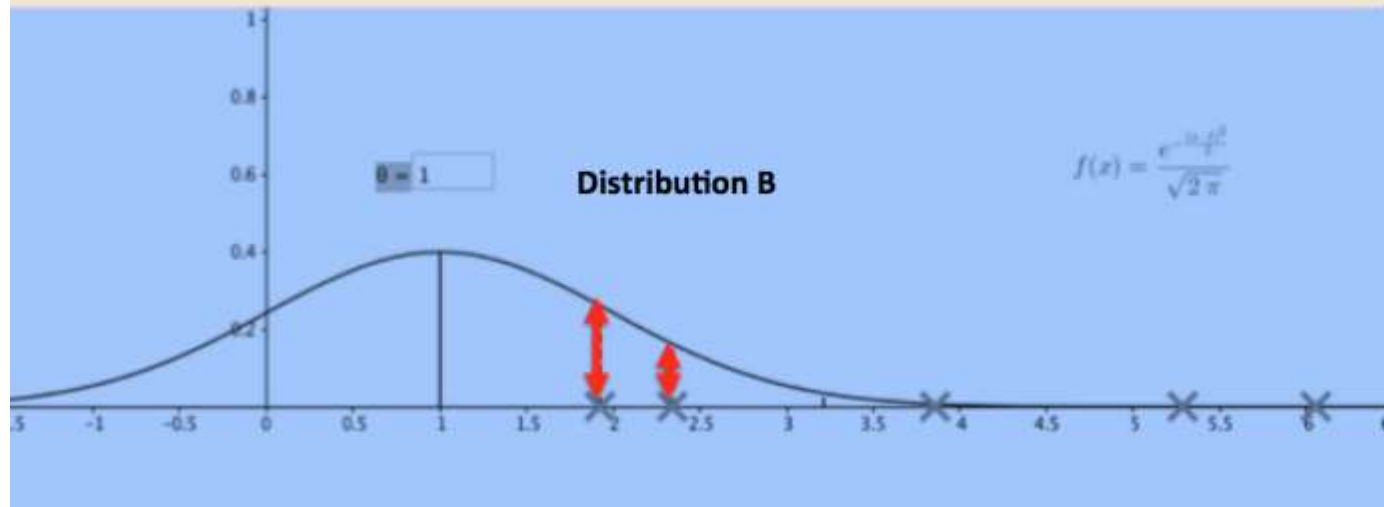
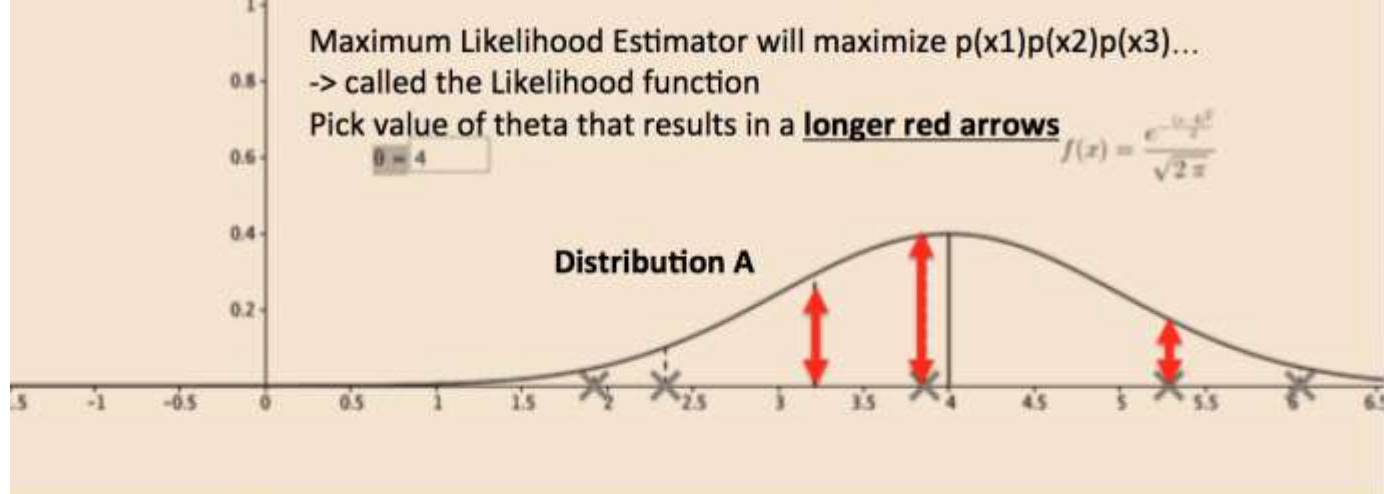
# Maximum Likelihood Estimation

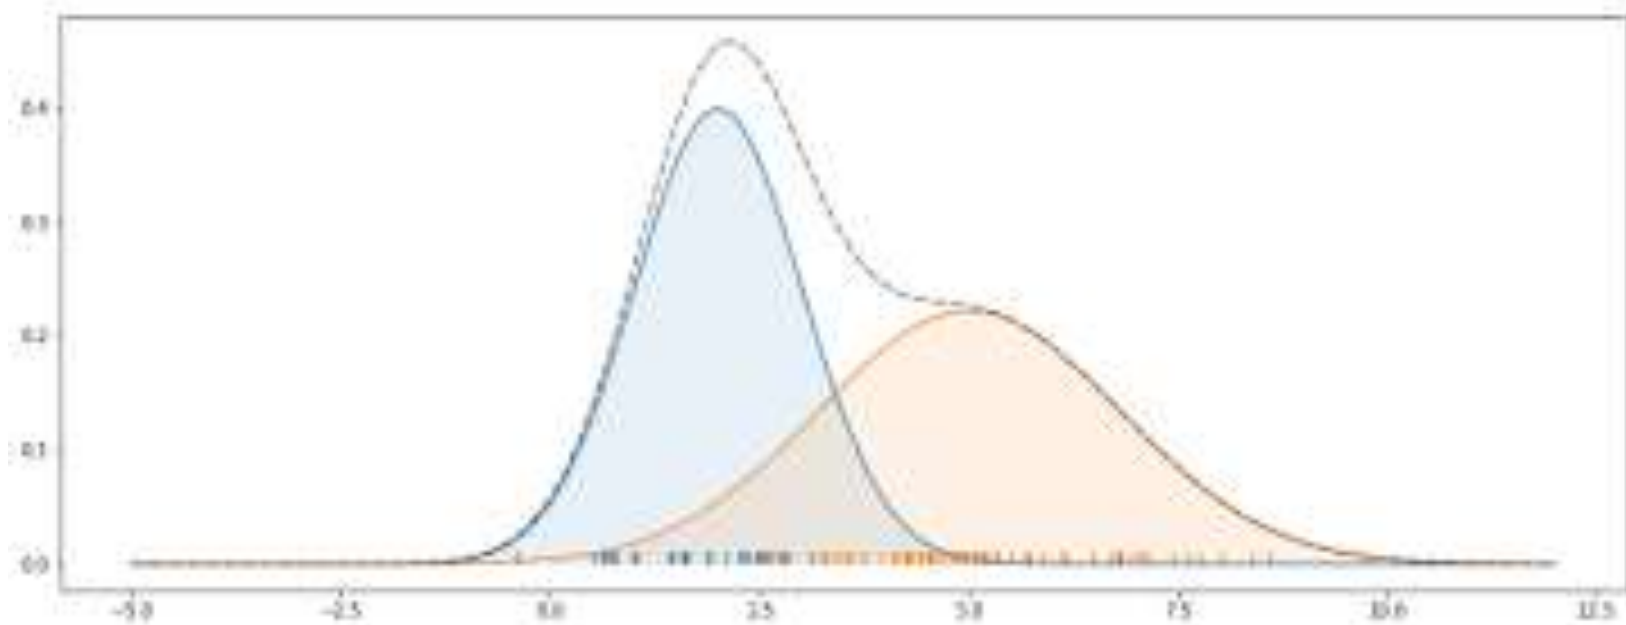
Maximum Likelihood Estimator will maximize  $p(x_1)p(x_2)p(x_3)\dots$

-> called the Likelihood function

Pick value of theta that results in a longer red arrows

$$f(x) = \frac{e^{-\frac{(x-\theta)^2}{2}}}{\sqrt{2\pi}}$$





# Mixture Models

- Data distribution  $p(x)$  assumed to be a **weighted sum** of  $K$  distributions

$$p(x) = \sum_{k=1}^K \pi_k p(x|\theta_k)$$

where  $\pi_k$ 's are the **mixing weights**:  $\sum_{k=1}^K \pi_k = 1$ ,  $\pi_k \geq 0$  (intuitively,  $\pi_k$  is the proportion of data generated by the  $k$ -th distribution)

- Each component distribution  $p(x|\theta_k)$  represents a “cluster” in the data

# Mixture Models

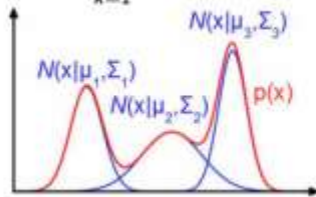
- Data distribution  $p(x)$  assumed to be a **weighted sum** of  $K$  distributions

$$p(x) = \sum_{k=1}^K \pi_k p(x|\theta_k)$$

where  $\pi_k$ 's are the **mixing weights**:  $\sum_{k=1}^K \pi_k = 1$ ,  $\pi_k \geq 0$  (intuitively,  $\pi_k$  is the proportion of data generated by the  $k$ -th distribution)

- Each component distribution  $p(x|\theta_k)$  represents a “cluster” in the data
- **Gaussian Mixture Model (GMM)**: component distributions are Gaussians

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$$





# Mixture Models

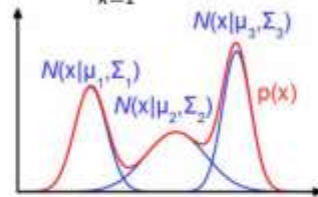
- Data distribution  $p(x)$  assumed to be a **weighted sum** of  $K$  distributions

$$p(x) = \sum_{k=1}^K \pi_k p(x|\theta_k)$$

where  $\pi_k$ 's are the **mixing weights**:  $\sum_{k=1}^K \pi_k = 1$ ,  $\pi_k \geq 0$  (intuitively,  $\pi_k$  is the proportion of data generated by the  $k$ -th distribution)

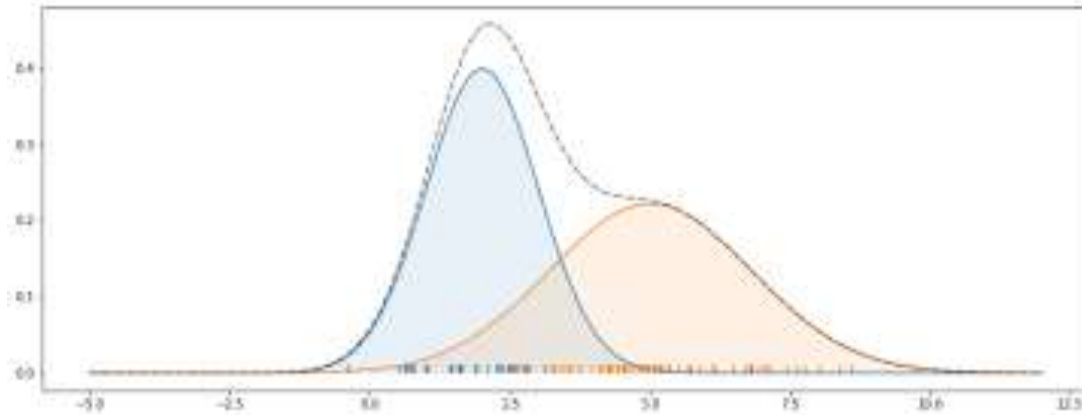
- Each component distribution  $p(x|\theta_k)$  represents a “cluster” in the data
- **Gaussian Mixture Model (GMM)**: component distributions are Gaussians

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$$



- Mixture models used in many data modeling problems, e.g.,
  - Unsupervised Learning: **Clustering (+density estimation)**
  - Supervised Learning: **Mixture of Experts** models

# Mixture Models



A GMM represents a **distribution** as

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k)$$

with  $\pi_k$  the **mixing coefficients**, where:

$$\sum_{k=1}^K \pi_k = 1 \quad \text{and} \quad \pi_k \geq 0 \quad \forall k$$

# Mixture Models

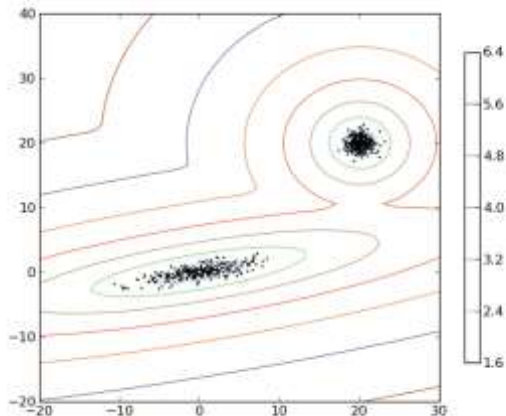
Most common mixture model: Gaussian mixture model (GMM)

- A GMM represents a **distribution** as

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k)$$

with  $\pi_k$  the mixing coefficients, where:

$$\sum_{k=1}^K \pi_k = 1 \quad \text{and} \quad \pi_k \geq 0 \quad \forall k$$



[http://scikit-learn.sourceforge.net/0.5/auto\\_examples/gmm/plot\\_gmm\\_pdf.html](http://scikit-learn.sourceforge.net/0.5/auto_examples/gmm/plot_gmm_pdf.html)

- Can think of the data  $\{\mathbf{x}_1, \mathbf{x}_n, \dots, \mathbf{x}_N\}$  using a “generative story”
  - For each example  $\mathbf{x}_n$ , first choose its cluster assignment  $z_n \in \{1, 2, \dots, K\}$  as

$$z_n \sim \text{Multinoulli}(\pi_1, \pi_2, \dots, \pi_K) \quad \text{aka “categorical”}$$

- Now generate  $\mathbf{x}$  from the Gaussian with id  $z_n$

$$\mathbf{x}_n | z_n \sim \mathcal{N}(\boldsymbol{\mu}_{z_n}, \boldsymbol{\Sigma}_{z_n})$$

# Resources

- Textbook
  - PRML (Bishop) – Chapter 9: 9.1,9.2,9.3.2
  - Pattern Classification (Duda, Hart, Stork)
    - 10.4.3,10.6.1,10.7.1,10.7.2,10.8,10.10
- Videos
  - [https://www.youtube.com/watch?v=REypj2sy\\_5U&list=PLBv09BD7ez\\_4e9LtmK626Evn1ion6ynrt](https://www.youtube.com/watch?v=REypj2sy_5U&list=PLBv09BD7ez_4e9LtmK626Evn1ion6ynrt)
  - <https://www.youtube.com/watch?v=rVfZHWTwXSA>
- Blog posts/Lecture Notes
  - [https://www.cse.iitk.ac.in/users/piyush/courses/pml\\_winter16/slides lec7.pdf](https://www.cse.iitk.ac.in/users/piyush/courses/pml_winter16/slides lec7.pdf)
  - <https://see.stanford.edu/materials/aimlcs229/cs229-notes8.pdf>
  - [https://www.cs.toronto.edu/~jlucas/teaching/csc411/lectures/lec15\\_16\\_handout.pdf](https://www.cs.toronto.edu/~jlucas/teaching/csc411/lectures/lec15_16_handout.pdf)
  - [https://mbernste.github.io/posts/gmm\\_em/](https://mbernste.github.io/posts/gmm_em/)
  - <https://www.ritchievink.com/blog/2019/05/24/algorithm-breakdown-expectation-maximization/>