

10.09.2024

Statistical Methods in AI (CS7.403)

Lecture-11: PCA – 2, PPCA, FA, Feature Selection

Ravi Kiran (ravi.kiran@iiit.ac.in)

<https://ravika.github.io>

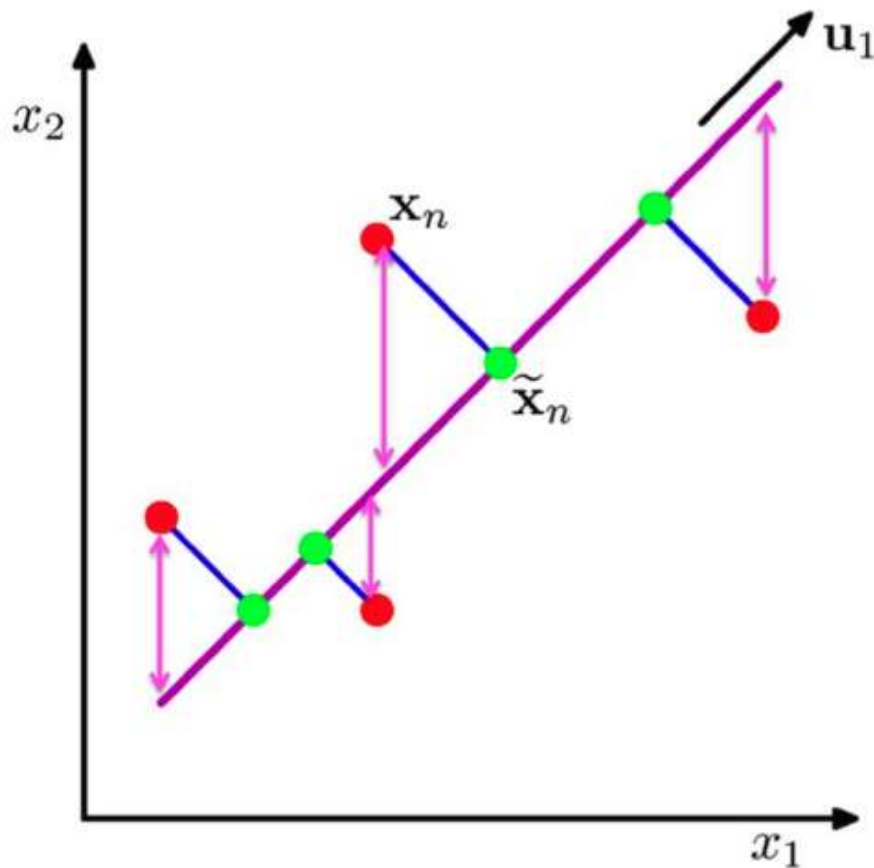


Center for Visual Information Technology (CVIT)

IIIT Hyderabad

PCA: Why eigenvectors ?

PCA vs linear regression



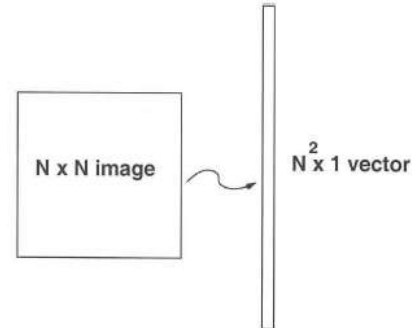
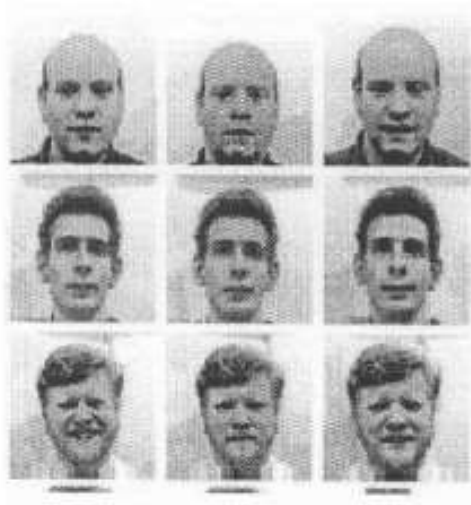
In contrast: in regression we'd minimize square error on *one* dimension (x_2) using a linear combination the *other dimensions*

Application to Faces

- Computation of low-dimensional basis (i.e., eigenfaces):

Step 1: obtain face images I_1, I_2, \dots, I_M (training faces)

(**very important:** the face images must be *centered* and of the same *size*)



Step 2: represent every image I_i as a vector Γ_i

Application to Faces

- Computation of the eigenfaces – cont.

Step 3: compute the average face vector Ψ :

$$\Psi = \frac{1}{M} \sum_{i=1}^M \Gamma_i$$

Step 4: subtract the mean face:

$$\Phi_i = \Gamma_i - \Psi$$

Step 5: compute the covariance matrix C :

$$C = \frac{1}{M} \sum_{n=1}^M \Phi_n \Phi_n^T = \frac{1}{M} A A^T \quad (N^2 \times N^2 \text{ matrix})$$

$$\text{where } A = [\Phi_1 \ \Phi_2 \ \cdots \ \Phi_M] \quad (N^2 \times M \text{ matrix})$$

Application to Faces

- Computation of the eigenfaces – cont.

Step 6: compute the eigenvectors u_i of $AA^T \rightarrow AA^T u_i = \lambda_i u_i$

The matrix AA^T is very large --> not practical !!

Application to Faces

- Computation of the eigenfaces – cont.

Step 6: compute the eigenvectors u_i of $AA^T \rightarrow AA^T u_i = \lambda_i u_i$

The matrix AA^T is very large --> not practical !!

Step 6.1: consider the matrix $A^T A$ ($M \times M$ matrix)

Step 6.2: compute the eigenvectors v_i of $A^T A$

$$A^T A v_i = \mu_i v_i$$

What is the relationship between u_i and v_i ?

Application to Faces

- Computation of the eigenfaces – cont.

Step 6: compute the eigenvectors u_i of $AA^T \rightarrow AA^T u_i = \lambda_i u_i$

The matrix AA^T is very large --> not practical !!

Step 6.1: consider the matrix $A^T A$ ($M \times M$ matrix)

Step 6.2: compute the eigenvectors v_i of $A^T A$

$$A^T A v_i = \mu_i v_i$$

What is the relationship between u_i and v_i ?

$$A^T A v_i = \mu_i v_i \Rightarrow AA^T A v_i = \mu_i A v_i \Rightarrow$$

$$u_i = A v_i \quad \text{and} \quad \lambda_i = \mu_i$$

Thus, AA^T and $A^T A$ have the same eigenvalues and their eigenvectors are related as follows: $u_i = A v_i$!!

Application to Faces

- Computation of the eigenfaces – cont.

Note 1: AA^T can have up to N^2 eigenvalues and eigenvectors.

Note 2: $A^T A$ can have up to M eigenvalues and eigenvectors.

Note 3: The M eigenvalues of $A^T A$ (along with their corresponding eigenvectors) correspond to the M *largest* eigenvalues of AA^T (along with their corresponding eigenvectors).

Step 6.3: compute the M best eigenvectors of AA^T : $u_i = Av_i$

(**important:** normalize u_i such that $\|u_i\| = 1$)

Step 7: keep only K eigenvectors (corresponding to the K largest eigenvalues)

Eigendecomposition vs. Singular Value Decomposition

- Eigendecomposition
 - Must be a diagonalizable matrix
 - Must be a square matrix
 - Matrix (n x n size) must have n linearly independent eigenvector
 - e.g. symmetric matrix ..

$$\boxed{A} = \boxed{P} \boxed{\Lambda} \boxed{P^{-1}}$$

Eigendecomposition vs. Singular Value Decomposition

- Eigendecomposition

- Must be a diagonalizable matrix
- Must be a square matrix
- Matrix ($n \times n$ size) must have n linearly independent eigenvector
 - e.g. symmetric matrix ..

$$\boxed{A} = \boxed{P} \boxed{\Lambda} \boxed{P^{-1}}$$

- Singular Value Decomposition

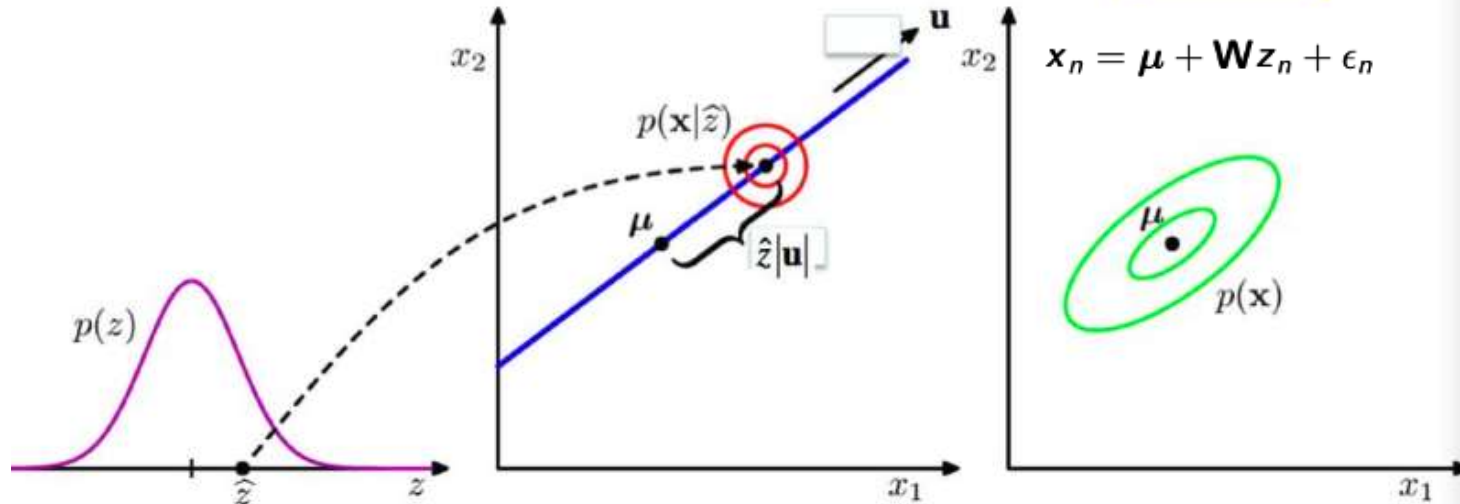
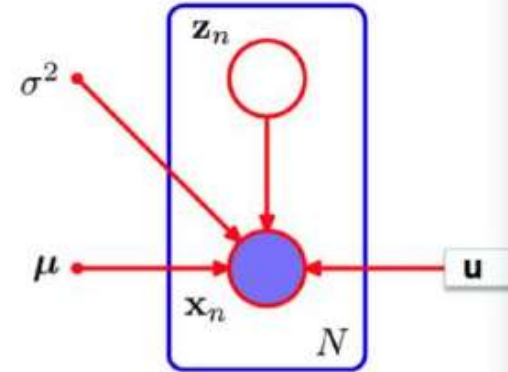
- Computable for any size ($M \times n$) of matrix

$$\boxed{A} = \boxed{U} \boxed{\Sigma} \boxed{V^T}$$

Probabilistic PCA

PCA

- Pick a *continuous* value z , which will be used to combine the “prototypes” \mathbf{u} in the model
- Pick the point \mathbf{x} from a spherical Gaussian centered on $z\mathbf{u}$



Probabilistic PCA

- Assume the following generative model for each observation x_n

$$x_n = \mathbf{W}z_n + \epsilon_n$$

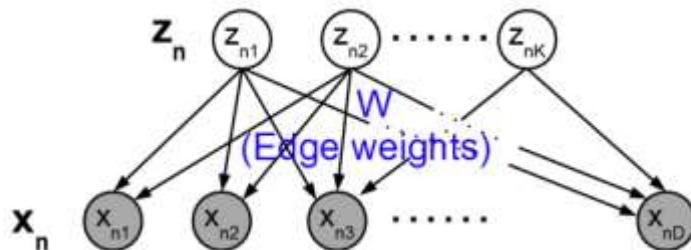
- Note: We'll assume data to be centered, otherwise $x_n = \mu + \mathbf{W}z_n + \epsilon_n$

Probabilistic PCA

- Assume the following generative model for each observation \mathbf{x}_n

$$\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$$

- Note: We'll assume data to be centered, otherwise $\mathbf{x}_n = \boldsymbol{\mu} + \mathbf{W}\mathbf{z}_n + \epsilon_n$
- Think of it as low dimensional $\mathbf{z}_n \in \mathbb{R}^K$ "generating" a higher-dimensional $\mathbf{x}_n \in \mathbb{R}^D$ via a mapping matrix $\mathbf{W} \in \mathbb{R}^{D \times K}$, plus some noise $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$

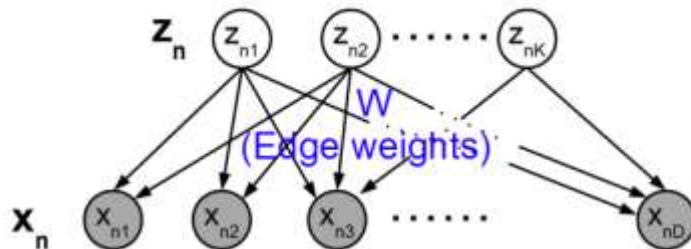


Probabilistic PCA

- Assume the following generative model for each observation x_n

$$x_n = \mathbf{W}z_n + \epsilon_n$$

- Note: We'll assume data to be centered, otherwise $x_n = \mu + \mathbf{W}z_n + \epsilon_n$
- Think of it as low dimensional $z_n \in \mathbb{R}^K$ "generating" a higher-dimensional $x_n \in \mathbb{R}^D$ via a mapping matrix $\mathbf{W} \in \mathbb{R}^{D \times K}$, plus some noise $\epsilon_n \sim \mathcal{N}(0, \sigma^2 \mathbf{I}_D)$



- Intuitively, this generative model is "inverse" of what the traditional PCA does. Here we assume a latent low-dim z_n that "generates" the high-dim x_n via the mapping \mathbf{W} (plus adding some noise)
- A directed graphical model linking z_n and x_n via "edge weights" \mathbf{W}

Probabilistic PCA

- Can also write $\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$ as each example \mathbf{x}_n being a linear comb. of columns of $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$, plus some example-specific random noise ϵ_n

$$\mathbf{x}_n = \sum_{k=1}^K \mathbf{w}_k z_{nk} + \epsilon_n$$

- The K columns of \mathbf{W} (each \mathbb{R}^D) are like “prototype vectors” shared by all examples. Each \mathbf{x}_n is a linear combination of these vectors (the combination coefficients are given by $\mathbf{z}_n \in \mathbb{R}^K$ which is basically the low-dim rep. of \mathbf{x}_n).

Probabilistic PCA

- Can also write $\mathbf{x}_n = \mathbf{W}\mathbf{z}_n + \epsilon_n$ as each example \mathbf{x}_n being a linear comb. of columns of $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_K]$, plus some example-specific random noise ϵ_n

$$\mathbf{x}_n = \sum_{k=1}^K \mathbf{w}_k z_{nk} + \epsilon_n$$

- The K columns of \mathbf{W} (each \mathbb{R}^D) are like “prototype vectors” shared by all examples. Each \mathbf{x}_n is a linear combination of these vectors (the combination coefficients are given by $\mathbf{z}_n \in \mathbb{R}^K$ which is basically the low-dim rep. of \mathbf{x}_n).
- Some examples:
 - In case of images, columns of \mathbf{W} would correspond to “basis images”
 - In case of text documents, columns of \mathbf{W} (with non-negativity imposed on it) would correspond to “topics” in the corpus

Probabilistic PCA - EM

- Since noise $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$ is Gaussian, the conditional distrib. of \mathbf{x}_n

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

- Given a set of observations $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, the goal is to learn \mathbf{W} and the low-dim. representation of data, i.e., $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$

Probabilistic PCA - EM

- Since noise $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$ is Gaussian, the conditional distrib. of \mathbf{x}_n

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

- Given a set of observations $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, the goal is to learn \mathbf{W} and the low-dim. representation of data, i.e., $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$
- Assume a Gaussian prior on the low-dimensional latent representation, i.e.,

$$p(\mathbf{z}_n) = \mathcal{N}(0, \mathbf{I}_K)$$

- Observed data: $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, latent variable: $\mathbf{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_N\}$
- Parameters: \mathbf{W}, σ^2
- The complete data log -likelihood

$$\begin{aligned} \log p(\mathbf{X}, \mathbf{Z} | \mathbf{W}, \sigma^2) &= \log \prod_{n=1}^N p(\mathbf{x}_n, \mathbf{z}_n | \mathbf{W}, \sigma^2) = \log \prod_{n=1}^N p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) p(\mathbf{z}_n) \\ &= \sum_{n=1}^N \{ \log p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) + \log p(\mathbf{z}_n) \} \end{aligned}$$

Benefits of Probabilistic PCA

- Can handle missing data (can treat it as latent variable in E step)
- Doesn't require computing the $D \times D$ cov. matrix of data and doing expensive eigen-decomposition. When K is small (i.e., we only want few eigen vectors), this is especially nice because only inverting $K \times K$ is required

Benefits of Probabilistic PCA

- Can handle missing data (can treat it as latent variable in E step)
- Doesn't require computing the $D \times D$ cov. matrix of data and doing expensive eigen-decomposition. When K is small (i.e., we only want few eigen vectors), this is especially nice because only inverting $K \times K$ is required
- Easy to “plug-in” PPCA as part of more complex problems, e.g., mixtures of PPCA models for doing nonlinear dimensionality reduction, or subspace clustering (i.e., clustering when data in each cluster lives on a lower dimensional subspace).
- Possible to give it a fully Bayesian treatment (which has many benefits such as inferring K)

Factor Analysis

- Similar to PPCA except that the Gaussian conditional distribution $p(\mathbf{x}_n | \mathbf{z}_n)$ has diagonal instead of spherical covariance

$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$$

where $\mathbf{\Psi}$ is a diagonal matrix

Probabilistic PCA - EM

- Since noise $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$ is Gaussian, the conditional distrib. of \mathbf{x}_n

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

Factor Analysis

- Similar to PPCA except that the Gaussian conditional distribution $p(\mathbf{x}_n | \mathbf{z}_n)$ has diagonal instead of spherical covariance

$$\mathbf{x}_n \sim \mathcal{N}(\mathbf{W}\mathbf{z}_n, \mathbf{\Psi})$$

where $\mathbf{\Psi}$ is a diagonal matrix

- In Factor Analysis, the projection matrix \mathbf{W} is also called the **Factor Loading Matrix** and \mathbf{z}_n is called the **factor scores** for example n

Probabilistic PCA - EM

- Since noise $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$ is Gaussian, the conditional distrib. of \mathbf{x}_n

$$p(\mathbf{x}_n | \mathbf{z}_n, \mathbf{W}, \sigma^2) = \mathcal{N}(\mathbf{W}\mathbf{z}_n, \sigma^2 \mathbf{I}_D)$$

PCA vs mixtures of Gaussians

Mixture of Gaussians

For each point:

- Pick the index of the (latent) Gaussian $Z=k$
- Pick the point \mathbf{x} from that the k -th Gaussian, $\mathbf{x} \sim \mathcal{N}(\mu_k, \Sigma_k)$

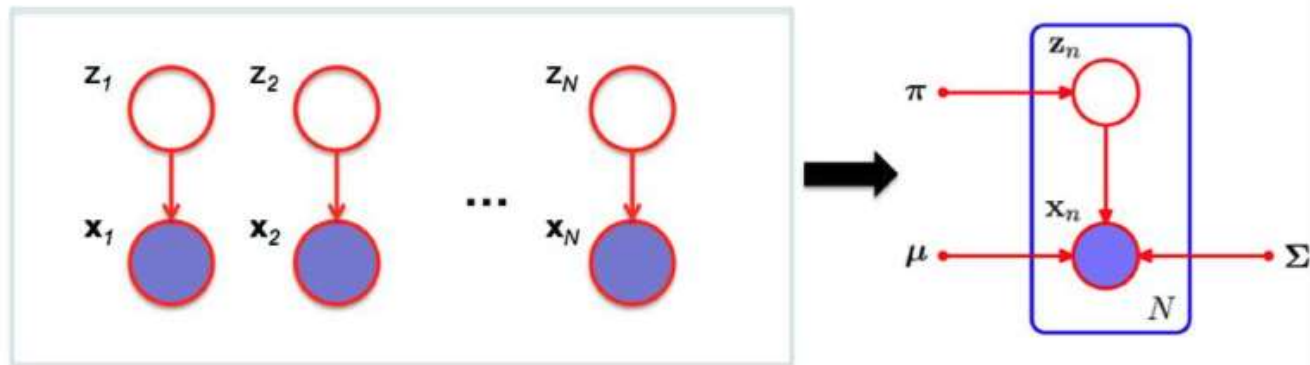
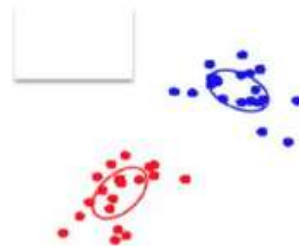
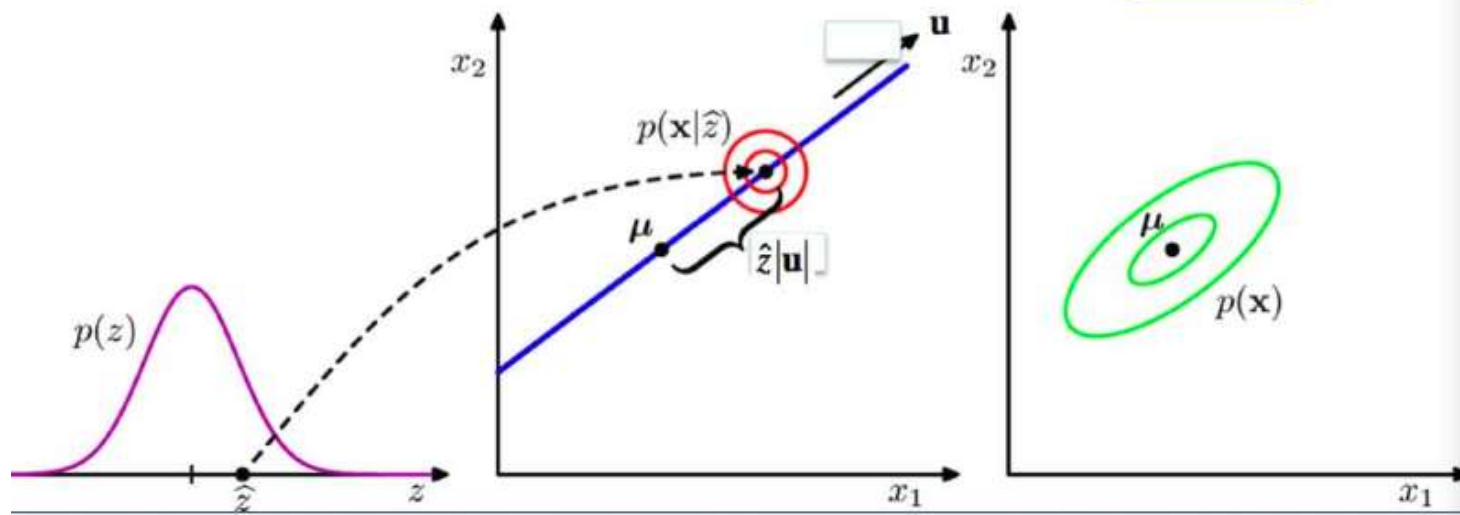
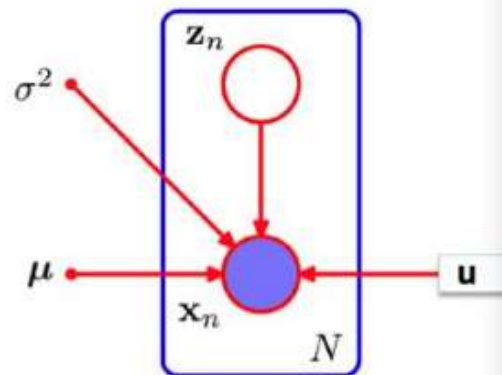


Plate notation

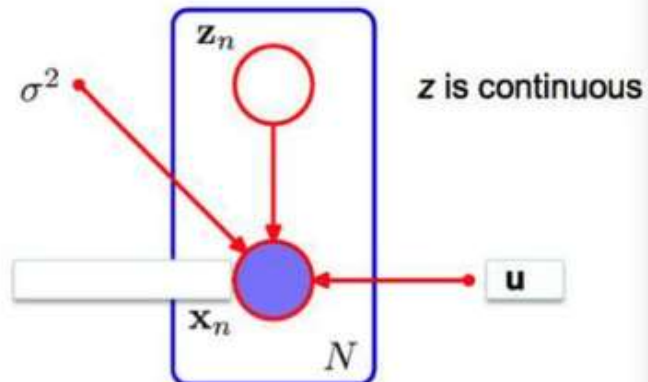
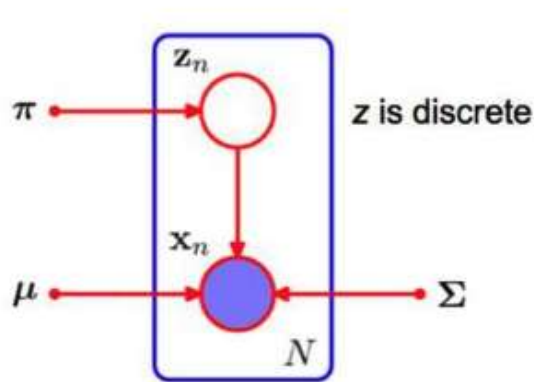
PCA vs mixtures of Gaussians

PCA

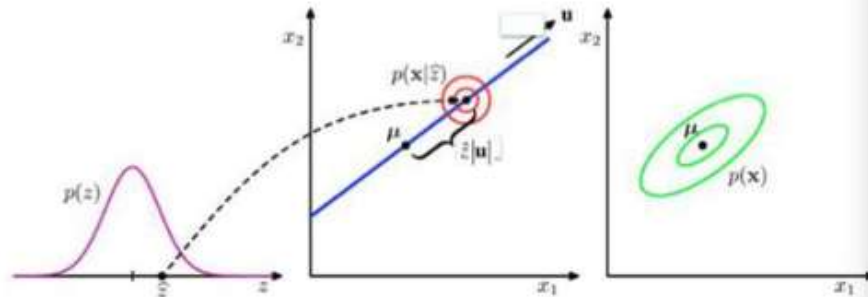
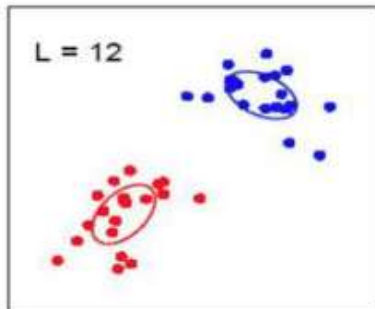
- Pick a *continuous* value z , which will be used to combine the “prototypes” \mathbf{u} in the model
- Pick the point \mathbf{x} from a spherical Gaussian centered on $z\mathbf{u}$



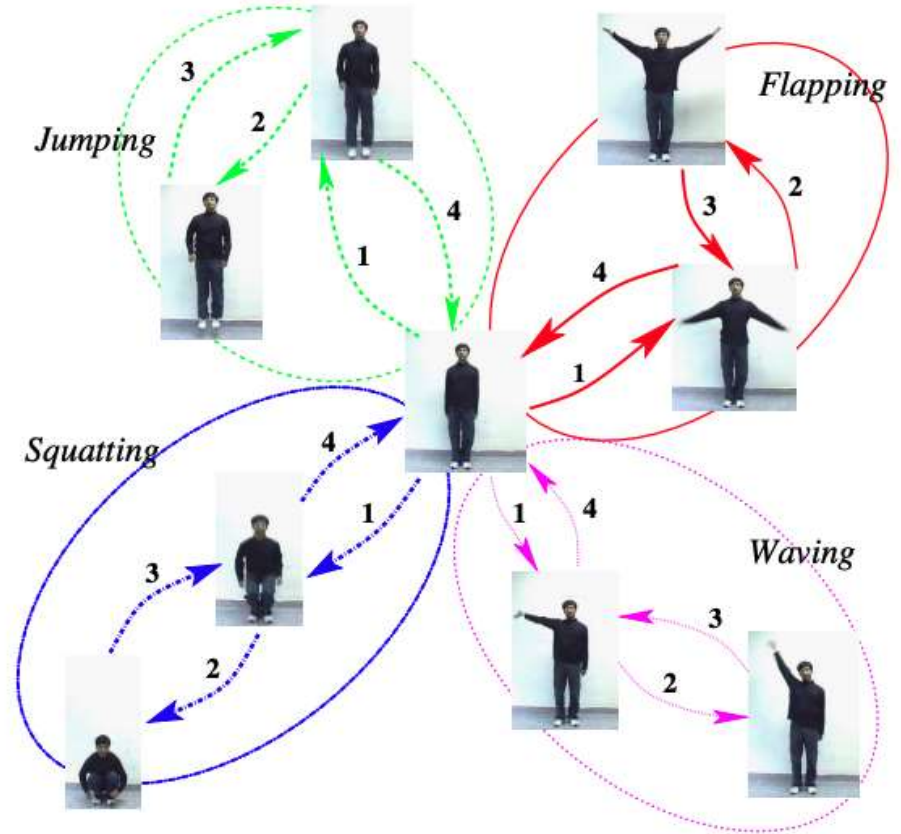
PCA vs mixtures of Gaussians



Comment: we can preprocess the data so that the mean is $\mathbf{0}$ to simplify the model

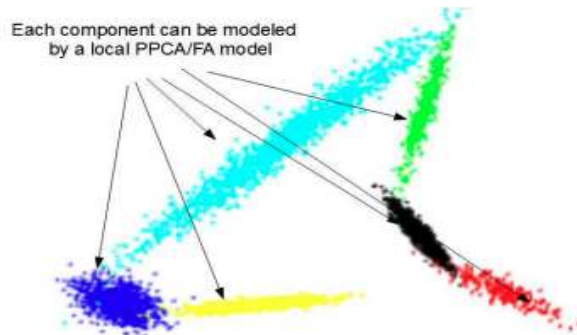


Recognizing Human Activities from Constituent Actions



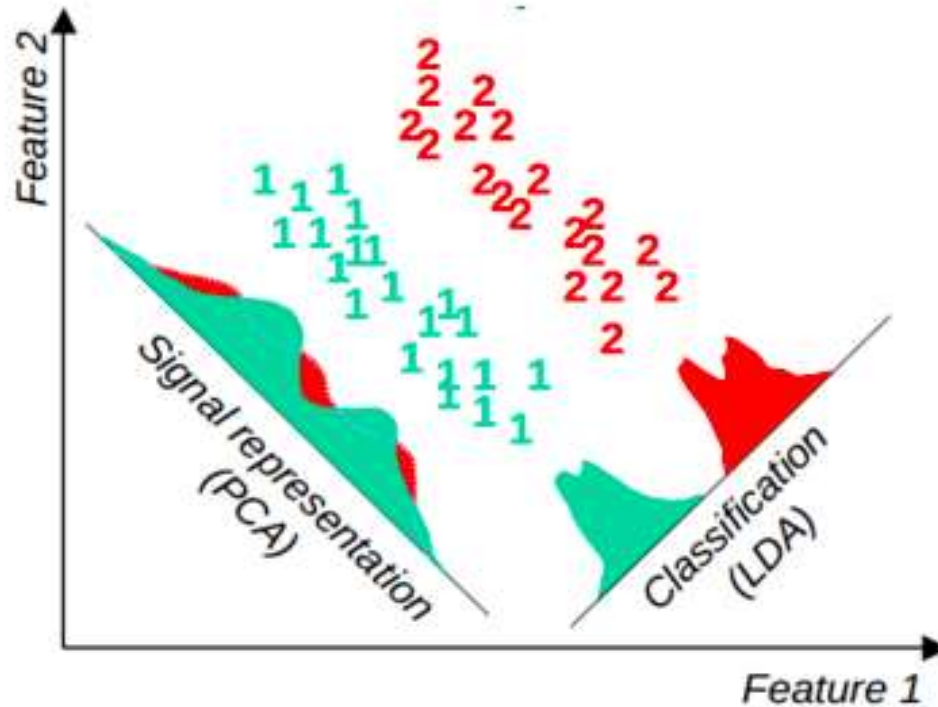
Mixture of PPCAs/Mixture of Factor Analyzers

- PPCA and FA learn a linear projection of the data (i.e., are linear dimensionality reduction methods)



- Similar to mixture of Gaussians, except that now each Gaussian is replaced by a PPCA or FA model

Linear Discriminant Analysis (LDA)



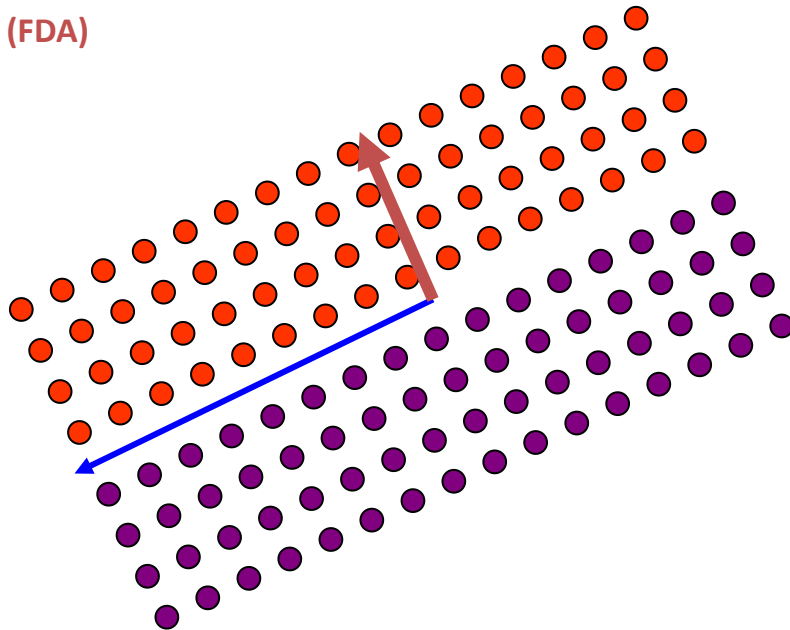
Linear Discriminant Analysis (LDA)

- Also known as “Fisher Discriminant”
- Does dimensionality Reduction
 - Also use the label “y”
 - Or Supervised Dimensionality Reduction
- PCA, LDA are both linear -- there are also **nonlinear** Dimensionality Reduction schemes (kernel PCA, t-SNE, autoencoders)

Beyond PCA

Are the maximal variance dimensions the relevant dimensions for preservation?

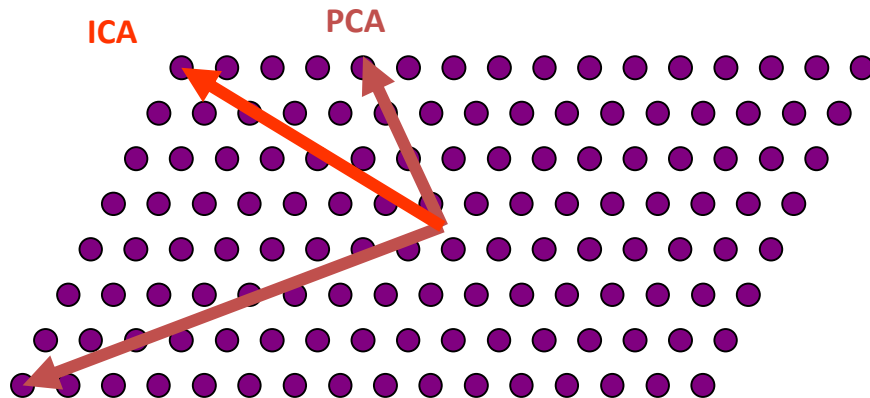
- Neighborhood Component Analysis (NCA)
- Fisher Discriminant analysis (FDA)



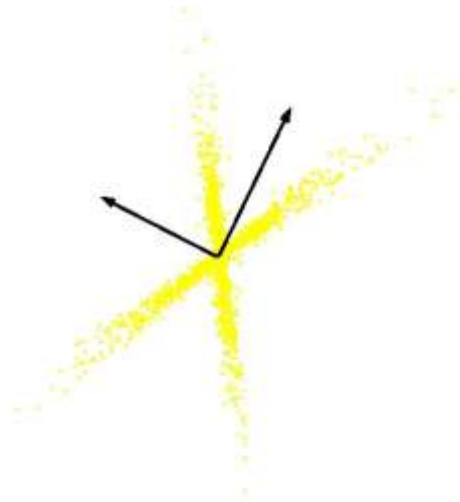
Beyond PCA

Should the goal be finding independent rather than pair-wise uncorrelated dimensions

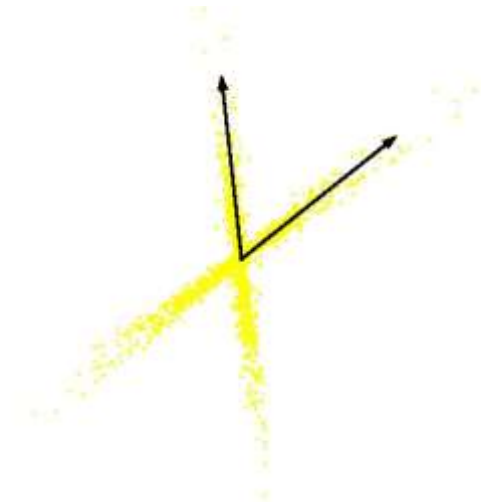
- Independent Component Analysis (ICA)



PCA vs ICA



PCA
(orthogonal coordinate)



ICA
(non-orthogonal coordinate)

Feature Selection Techniques

Supervised Feature Selection

Unsupervised Feature Selection

Filters method

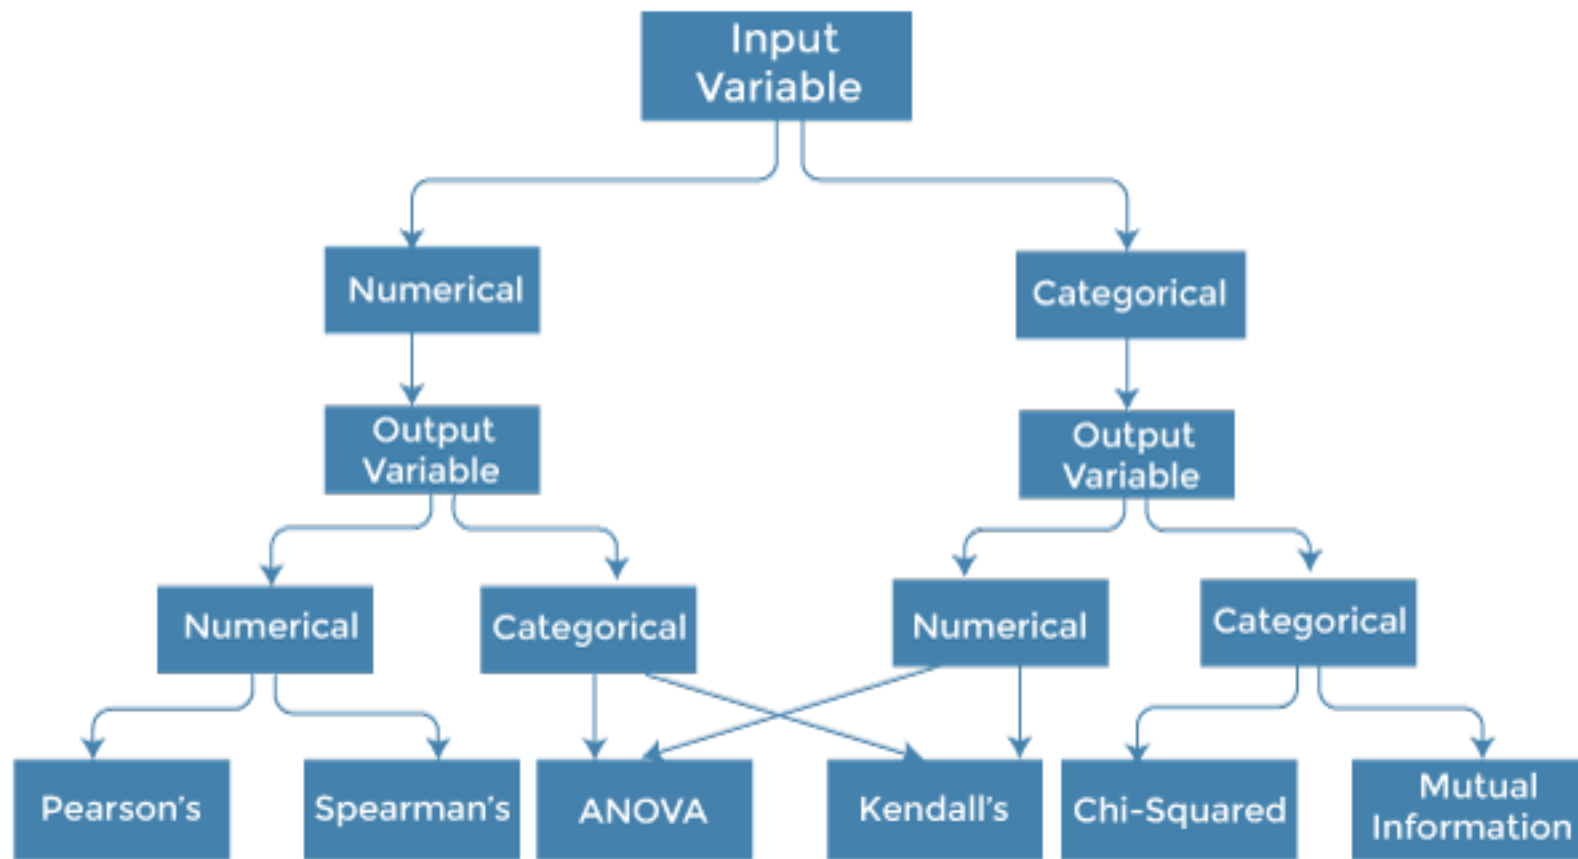
- Missing value
- Information gain
- Chi-square Test
- Fisher's Score

Embedded method

- Regularization L1, L2

Wrappers method

- Forward Feature Selection
- Backward Feature Selection
- Exhaustive Feature Selection
- Recursive Feature Elimination



References

- [Bishop PRML, 12.1, 12.2, 12.4](#)
- https://jeremy9959.net/Math-3094-UConn/published_notes/notes/PCA.pdf
- https://www.cse.iitk.ac.in/users/piyush/courses/pml_winter16/slides_lec10.pdf